



Radeon Voice Skill Foundry

SPEAK THE SOP. PROVE THE SKILL.

A local Agent can observe clicks.

It cannot infer the hidden rules.

Voice captures exceptions and prohibited actions. Server-held evidence checks level, clipping, silence, source integrity, and the original ASR transcript before promotion.

29.42%

fewer output tokens

compact structured output

30.03%

lower generation latency

same W7900 + same model

100/100

voice evidence gate

20.39 s SOP WAV passed

VOICE-TO-VERIFIED-SKILL

1

Speak + demonstrate

Private SOP and aligned action trace

2

Measure the source

Level, clipping, silence, audio hash

3

Compile locally

Qwen3-ASR + Qwen3-4B on Radeon

4

Prove + reuse

Policy, tests, receipts, skill memory

PROOF-CARRYING OUTPUT

- ☒ GAIA-compatible SKILL.md
- ☒ least-privilege policy.yaml
- ☒ positive + adversarial fixtures
- ☒ **governance receipts**
- ☒ hash-bound proof_bundle.json
- ☒ source-bound voice_evidence.json
- ☒ versioned procedural memory

Radeon Pro W7900-class | gfx1100 | ROCm 7.2.1

Qwen3-ASR-0.6B -> Qwen3-4B-Instruct -> deterministic verifier

mail.send = DENY

Local audio proof: 7/7

Track 2 | Team N/A | github.com/Chengyuann/radeon-voice-skill-foundry

Actions capture what happened. Voice captures why, when, and what must never happen.