



# Radeon Voice Skill Foundry

CROSS-MODAL POLICY INDUCTION

**A local Agent can observe clicks.**

**It cannot infer the hidden rules.**

Speech explains why, when, and what must never happen. Demonstrated actions show tools and parameters. Local policy provides authority. The W7900 compiles the three into a least-privilege skill and proves it before promotion.

**257.65**

**output tokens/s**

vLLM graph at concurrency 8

**12.47x**

**serving uplift**

versus serialized Transformers

**85.35x**

**ASR aggregate RT**

native Qwen3-ASR batch 8

**PASS**

**voice evidence**

internal deterministic gate

**7/7**

**adversarial proof**

mail.send remains DENY

**68/68**

**regression tests**

typecheck + production build

SPEECH + DEMONSTRATION + LOCAL POLICY

1

**Hear hidden rules**

exceptions, conditions and prohibitions

2

**Observe real tools**

action trace, parameters and state

3

**Retrieve authority**

local policy evidence and boundaries

4

**Prove + remember**

7/7 + hashed receipts + versioned memory

VERIFIED SKILL PACKAGE

- ✓ portable Agent Skill Markdown + policy
- ✓ positive and adversarial fixtures
- ✓ hashed receipts + artifact integrity fields
- ✓ source-bound Voice Evidence
- ✓ versioned memory + child-run lineage

LIVE PRODUCT PATH

- 01 **Cloudflare Pages module UI**
- 02 **authenticated same-origin gateway**
- 03 **W7900 Qwen3-ASR + Qwen3-4B**
- 04 **deterministic safety + proof kernel**

**mail.send = DENY**

**7/7 PASS**

WATCH THE REAL PATH

**radeon-voice-skill-foundry.pages.dev**

Product Demo: public W7900 workflow | Technical evidence: GitHub Release + evidence index

Track 2 | Chengyuan Ma | [github.com/Chengyuann/radeon-voice-skill-foundry](https://github.com/Chengyuann/radeon-voice-skill-foundry)

Not speech-to-text-to-chat: cross-modal evidence becomes a verified Agent Skill.