



Radeon Voice Skill Foundry

SPEAK THE SOP. PROVE THE SKILL.

A local Agent can observe clicks.

It cannot infer the hidden rules.

Voice captures exceptions and prohibited actions. The W7900 compiles those rules into a least-privilege skill, then adversarial replay proves it before promotion.

257.65

output tokens/s

vLLM graph at concurrency 8

12.47x

serving uplift

versus serialized Transformers

85.35x

ASR aggregate RT

native Qwen3-ASR batch 8

PASS

voice evidence

internal deterministic gate

7/7

adversarial proof

mail.send remains DENY

63/63

regression tests

typecheck + production build

VOICE-TO-VERIFIED-SKILL

1

Speak + demonstrate

Private SOP and aligned action trace

2

Measure + transcribe

Voice quality evidence + local ASR

3

Compile on W7900

local retrieval + constraints + least privilege

4

Prove + remember

7/7 + hashed receipts + versioned memory

VERIFIED SKILL PACKAGE

- ✓ portable Agent Skill Markdown + policy
- ✓ positive and adversarial fixtures
- ✓ hashed receipts + artifact integrity fields
- ✓ source-bound Voice Evidence
- ✓ versioned memory + child-run lineage

LIVE PRODUCT PATH

- 01 **Cloudflare Pages module UI**
- 02 **authenticated same-origin gateway**
- 03 **W7900 Qwen3-ASR + Qwen3-4B**
- 04 **deterministic safety + proof kernel**

mail.send = DENY

7/7 PASS

WATCH THE REAL PATH

radeon-voice-skill-foundry.pages.dev

Product Demo: public W7900 workflow | Technical evidence: GitHub Release + evidence index

Track 2 | Team N/A | github.com/Chengyuann/radeon-voice-skill-foundry

Actions capture what happened. Voice captures why, when, and what must never happen.