



Radeon Voice Skill Foundry

CROSS-MODAL POLICY INDUCTION

A local Agent can observe clicks.

It cannot infer the hidden rules.

Speech explains why, when, and what must never happen. Demonstrated actions show tools and parameters. Local policy provides authority. The W7900 compiles the three into a least-privilege skill and proves it before promotion.

257.65

output tokens/s

vLLM graph at concurrency 8

12.47x

serving uplift

versus serialized Transformers

85.35x

ASR aggregate RT

native Qwen3-ASR batch 8

PASS

voice evidence

internal deterministic gate

7/7

adversarial proof

mail.send remains DENY

68/68

regression tests

typecheck + production build

SPEECH + DEMONSTRATION + LOCAL POLICY

1

Hear hidden rules

exceptions, conditions and prohibitions

2

Observe real tools

action trace, parameters and state

3

Retrieve authority

local policy evidence and boundaries

4

Prove + remember

7/7 + hashed receipts + versioned memory

VERIFIED SKILL PACKAGE

- ✓ portable Agent Skill Markdown + policy
- ✓ positive and adversarial fixtures
- ✓ hashed receipts + artifact integrity fields
- ✓ source-bound Voice Evidence
- ✓ versioned memory + child-run lineage

LIVE PRODUCT PATH

- 01 **Cloudflare Pages module UI**
- 02 **authenticated same-origin gateway**
- 03 **W7900 Qwen3-ASR + Qwen3-4B**
- 04 **deterministic safety + proof kernel**

mail.send = DENY

7/7 PASS

WATCH THE REAL PATH

radeon-voice-skill-foundry.pages.dev

Product Demo: public W7900 workflow | Technical evidence: GitHub Release + evidence index

Track 2 | Chengyuan Ma | github.com/Chengyuann/radeon-voice-skill-foundry

Not speech-to-text-to-chat: cross-modal evidence becomes a verified Agent Skill.